

# Localization and mapping

From EKF to SLAM

Roberto Masocco

`roberto.masocco@uniroma2.it`

University of Rome Tor Vergata

Department of Civil Engineering and Computer Science Engineering

June 12, 2024



**TOR VERGATA**  
UNIVERSITY OF ROME

School of Engineering

- **Simulink Wrapper Node**

- ▶ ROS 2 node that wraps a Simulink Model, similarly to the MARTe2 SimulinkWrapperGAM.
- ▶ Deploy and test on a Turtlebot with a simple trajectory tracking algorithm fully onboard.

- **Comau e.DO manipulator ROS 2 integration**

- ▶ Setting up SBC (Nvidia Jetson TX2).
- ▶ Setting up Docker container.
- ▶ Setting up basic/advanced ROS 2 software for robot usage.

- **Improving zed\_drivers**

- ▶ Simple driver for non-GPU platforms based only on OpenCV.
- ▶ Simple driver for non-GPU platforms based on Open Capture API from Stereolabs.
- ▶ Deploy and test on a robot for visual servoing.

# Roadmap

- 1 The perception problem
- 2 Common interfaces
- 3 The tf2 library
- 4 The Extended Kalman Filter
- 5 The mapping problem
- 6 Simultaneous Localization And Mapping

# Roadmap

- 1 The perception problem
- 2 Common interfaces
- 3 The tf2 library
- 4 The Extended Kalman Filter
- 5 The mapping problem
- 6 Simultaneous Localization And Mapping

# The perception problem

## Definition

To be able to operate autonomously, a robot must continuously answer the following questions:

- **Where am I?**
- **What is this place?**

Thus, it must be able to **perceive** the environment, gathering information useful to:

- **localize itself**, *i.e.*, continuously estimate its **pose** as **both position and orientation** in 3D space;
- **map the environment**, *i.e.*, build a **representation** of the environment useful to **navigate** within it.

# The perception problem

## Challenges

The perception problem is challenging because:

- the environment may be **partially observable**, *i.e.*, the robot can only perceive a **subset** of it, and need to update its information in real time;
- the environment may be **dynamic**, *i.e.*, it can change over time;
- measurements are always subject to **noise**.

The perception problem is usually solved by **sensor fusion**, *i.e.*, combining information from **multiple sensors** to obtain a more **accurate** and **reliable** estimate of the environment, possibly accounting for **sensor faults**.

# The perception problem

## Tools for the job

The tools that robots use to gather **measurements** from the environment are called **sensors**.

They can be classified as:

- **proprioceptive**, *i.e.*, measuring robotic interaction with the environment (e.g., **encoders**, **GPS**, **IMUs**);
- **exteroceptive**, *i.e.*, measuring the environment itself (e.g., **cameras**, **LiDARs**, **radars**);
- **interoceptive**, *i.e.*, measuring the robot's internal state.

# The perception problem

## Tools for the job

As any other measurement tool, sensors are based on **physical principles** and **energy exchanges**, translating the information they gather into **electrical signals** that can be acquired and/or processed by a computer.

They are usually characterized by at least:

- a **digital** or **analog encoding** of the measurement;
- a **frame of reference** in which the measurement is expressed;
- **accuracy** and **uncertainty** parameters.



# The perception problem

Tools for the job




Figure 1: Rotary encoder working principle.

# The perception problem

Tools for the job



Figure 2: GPS module for drones.

# The perception problem

Tools for the job




Figure 3: Inertial Measurement Unit (IMU).

# The perception problem

Tools for the job




Figure 4: Light Detection and Ranging (LiDAR) sensor.

# The perception problem

Tools for the job



Figure 5: ZED 2i stereo camera.

# Roadmap

- 1 The perception problem
- 2 Common interfaces**
- 3 The tf2 library
- 4 The Extended Kalman Filter
- 5 The mapping problem
- 6 Simultaneous Localization And Mapping

# Common interfaces

A standard ROS 2 installation offers many **interface packages** (*i.e.*, messages), to provide **standard data types** to communicate sensor measurements and related data.

The most important are:

- `sensor_msgs`, for **sensor measurements**;
- `geometry_msgs`, for **geometric data**;
- `nav_msgs`, for **navigation data**.

It is suggested to **always use these message types**, plus common best practices, to ensure full interoperability between sensor drivers and localization and mapping algorithms.

Try to `ros2 interface show` these messages to understand their structure!

# sensor\_msgs

Common interfaces for sensors

- Imu
- JointState
- CameraInfo and Image
- LaserScan
- PointCloud2
- Temperature
- NavSatFix
- Illuminance
- ...



- Vector3Stamped
- QuaternionStamped
- PoseWithCovarianceStamped
- TwistWithCovarianceStamped
- TransformStamped (used by tf2!)
- AccelWithCovarianceStamped
- ...

- Odometry (Header, body (child) frame ID, PoseWithCovariance, TwistWithCovariance)
- Path
- OccupancyGrid
- ...

# Roadmap

- 1 The perception problem
- 2 Common interfaces
- 3 The tf2 library**
- 4 The Extended Kalman Filter
- 5 The mapping problem
- 6 Simultaneous Localization And Mapping

# Rigid transformations

When a robot moves in space, it is important to keep track of its **position** and **orientation** with respect to a **reference frame**.

Sensors measuring this information, as well as many more, are **mounted** on the robot, in fixed positions and orientations.

To process these measurements, they must first be **transformed** from the **body frame** into a common reference frame, usually called:

- **world frame** (world origin), in the case of **global localization**;
- **local frame**, or **odom frame** (robot starting point), in the case of **local localization**.

Such **rigid transformations** are **isometries**. They must be applied to almost every sensor measurement, and are usually **composable**.

We would like the middleware to provide tools to do this almost automatically...

# The tf2 library

tf2 is the **standard ROS 2 library** to handle rigid transformations. It allows to:

- **broadcast** and **listen** to transformations;
- optimize the broadcasting of **static transformations** vs the buffering of the others;
- **transform** any kind of sensor data from one frame to another, making efficient computations in C++ code relying on the Eigen mathematical library;
- broadcast static **robot descriptions** from URDF files, listing links and joints and how they are connected, resulting in a **tree-like structure**;
- **command-line tools** to introspect the tf tree.




Figure 6: Example of robot description with tf2.

# The tf2 library




Figure 7: Broadcasted robot descriptions, plus real-time transformations given by localization systems.

# Roadmap

- 1 The perception problem
- 2 Common interfaces
- 3 The tf2 library
- 4 The Extended Kalman Filter**
- 5 The mapping problem
- 6 Simultaneous Localization And Mapping

# Odometric reconstruction

An elementary sensor fusion technique

Suppose you have a robot equipped with  $m$  **sensors** producing **measurements**  $q_i$ ,  $i = 1, \dots, m$ , and a **motion model** that predicts the **state** of the robot  $x_k$  at time  $k$  given the state  $x_{k-1}$  at time  $k - 1$  and the control input  $u_k$  applied to the robot.

You can use these data to perform an **odometric reconstruction**: integrating the motion model with sensor data and the **estimated state** of the robot.



# Odometric reconstruction

An elementary sensor fusion technique

Suppose you have a robot equipped with  $m$  **sensors** producing **measurements**  $q_i$ ,  $i = 1, \dots, m$ , and a **motion model** that predicts the **state** of the robot  $x_k$  at time  $k$  given the state  $x_{k-1}$  at time  $k - 1$  and the control input  $u_k$  applied to the robot.

You can use these data to perform an **odometric reconstruction**: integrating the motion model with sensor data and the **estimated state** of the robot.

**Measurements are affected by noise, that this technique does not account for.**

Noise may even be due to **physical phenomena** like slippage, vibrations, or sensor faults.

# Modelling noise

## Gaussian random variables

Suppose that  $m = 2$ . At any given time, the **measurement** of a quantity  $x$  can be modeled as a **Gaussian random variable** centered around the **true value** of  $x$  with a certain **standard deviation**:

- $q_1 \sim \mathcal{N}(x, \sigma_1^2)$ ;
- $q_2 \sim \mathcal{N}(x, \sigma_2^2)$ .

Having distinct sensors makes these variables **independent**.

These assumptions are reasonable because:

- of well-known results, e.g., the Central Limit Theorem;
- the **average** is the real value of  $x$  if we rule out **systematic errors** by, e.g., calibrating the sensor;
- the **variance** models the **dispersion** of the measurements around the central value.

# Modelling noise

## Gaussian random variables

Note that, by relying on **independence**, **Bayes' theorem**, and the **properties** of the Gaussian distribution:

- $q_1 \sim \mathcal{N}(x, \sigma_1^2) \Rightarrow x \sim \mathcal{N}(q_1, \sigma_1^2) ;$

- $q_2 \sim \mathcal{N}(x, \sigma_2^2) \Rightarrow x \sim \mathcal{N}(q_2, \sigma_2^2) ;$

*i.e.*,  $x$  **is also a Gaussian random variable** centered around a measurement.

# Modelling noise

## Combining variables

**A priori**, and in particular before getting any measurements, we can **assume that  $x$  has a uniform distribution**, making  $p(x)$  **independent** from both  $p(q_1)$  and  $p(q_2)$ .

Relying on the same previous results, we can **combine the two measurements**  $q_1$  and  $q_2$  to obtain a **better estimate** of  $x$ , as  $x \sim \mathcal{N}(q, \sigma^2)$ , where:

- $q = \frac{\sigma_2^2 q_1 + \sigma_1^2 q_2}{\sigma_1^2 + \sigma_2^2}$  ;
- $\sigma^2 = \frac{\sigma_1^2 \sigma_2^2}{\sigma_1^2 + \sigma_2^2}$  ;

*i.e.*,  $x$  is a new Gaussian random variable centered around a **weighted average** of the measurements, with a **variance** that decreases as the **uncertainty** of the measurements decreases.

# Modelling noise

## Combining variables

To see this, start from using **independence** to see what it means to get  $q_1$  and  $q_2$ :

$$p((q_1, q_2)|x) = p(q_1|x)p(q_2|x) = \dots = \eta e^{-\frac{(x-q)^2}{2\sigma^2}},$$

where

$$\eta = \frac{1}{2\pi\sigma_1\sigma_2} e^{\frac{q^2}{2\sigma^2} - \frac{q_1^2}{2\sigma_1^2} - \frac{q_2^2}{2\sigma_2^2}}.$$

Then, **Bayes** tells you that

$$p(x|(q_1, q_2)) = p((q_1, q_2)|x) \frac{p(x)}{p((q_1, q_2))} = \frac{\eta p(x)}{p((q_1, q_2))} e^{-\frac{(x-q)^2}{2\sigma^2}} = \tilde{\eta} e^{-\frac{(x-q)^2}{2\sigma^2}},$$

where  $\tilde{\eta}$  is constant due to all the hypotheses made, and in particular

$$\int_{-\infty}^{+\infty} p(x|(q_1, q_2)) dx = 1 \implies \tilde{\eta} = \frac{1}{\sqrt{2\pi\sigma^2}}.$$

# Combining two measurements

## Step by step

Suppose you have to update an **estimate**  $\hat{x}$  of a quantity  $x$  and its **variance**  $\Sigma^2$  with two measurements  $q_1$  and  $q_2$ :

① Measurement  $q_1$  arrives:  $\hat{x}_1 = q_1$ ,  $\Sigma_1^2 = \sigma_1^2$ .

② Measurement  $q_2$  arrives.

③ 
$$\hat{x}_{1,2} = \frac{\Sigma_1^2 q_2 + \sigma_2^2 \hat{x}_1}{\Sigma_1^2 + \sigma_2^2} = \hat{x}_1 + \frac{\Sigma_1^2}{\Sigma_1^2 + \sigma_2^2} (q_2 - \hat{x}_1).$$

④ 
$$\Sigma_{1,2}^2 = \frac{\Sigma_1^2 \sigma_2^2}{\Sigma_1^2 + \sigma_2^2} = \left(1 - \frac{\Sigma_1^2}{\Sigma_1^2 + \sigma_2^2}\right) \Sigma_1^2.$$

# Combining many measurements

## The first innovation

Suppose now that you can receive  $m$  measurements in **subsequent steps**. Then:

$$\hat{x}_{k+1} = \hat{x}_k + V_{k+1}(q_{k+1} - \hat{x}_k), \quad (1)$$

$$\Sigma_{k+1}^2 = (1 - V_{k+1})\Sigma_k^2, \quad (2)$$

where:

$$V_{k+1} = \frac{\Sigma_k^2}{\Sigma_k^2 + \sigma_{k+1}^2}, \quad (3)$$

and  $V_{k+1} \in (0, 1)$ . It is a **sampled approximation**.

Note what happens when the new  $k + 1$ -th measurement is either **very accurate** ( $\sigma_{k+1}^2 \ll \Sigma_k^2$ ) or **very noisy** ( $\sigma_{k+1}^2 \gg \Sigma_k^2$ ).

# Kalman filter

## The scalar case

Suppose you have a **dynamic system** with a **discrete-time model**:

$$x_{k+1} = ax_k + bu_k + \omega_k, \quad (4)$$

with  $\omega_k \sim \mathcal{N}(0, \sigma_\omega^2)$ , and **measurements that depend on the state**:

$$z_k = cx_k + \nu_k, \quad (5)$$

where  $\nu_k \sim \mathcal{N}(0, \sigma_\nu^2)$  is the **noise**.



# Kalman filter

## The scalar case

Initialize the **estimate** as  $\hat{x}_0$  (0 is mostly fine) and the **covariance** as  $P_0$  (a rough estimate of the uncertainty).

Then, relying on the same properties as the previous cases:

$$\hat{x}_{k+1}^- = a\hat{x}_k + bu_k, \quad (6a)$$

$$P_{k+1}^- = a^2 P_k + \sigma_\omega^2, \quad (6b)$$

$$K_{k+1} = \frac{P_{k+1}^- c}{c^2 P_{k+1}^- + \sigma_\nu^2}, \quad (6c)$$

$$\hat{x}_{k+1} = \hat{x}_{k+1}^- + K_{k+1}(z_{k+1} - c\hat{x}_{k+1}^-), \quad (6d)$$

$$P_{k+1} = (1 - K_{k+1}c)P_{k+1}^-. \quad (6e)$$

(6a)-(6b) compose the **prediction step** of the **a priori estimate**, (6c) is the **Kalman gain**, and (6d)-(6e) are the **correction step**, applying the **innovation term** to compute the **a posteriori estimate**.

# Kalman filter

## The vectorial case

The **vectorial case** is a **generalization** of the scalar case, using **vectors** and **covariance matrices**.

Suppose that now you have this model:

$$x_{k+1} = Ax_k + Bu_k + \omega_k, \quad (7)$$

where  $A$  and  $B$  are **matrices**,  $\omega_k \sim \mathcal{N}(0, Q_\omega)$ , and the **measurement** is:

$$z_k = Cx_k + \nu_k, \quad (8)$$

where  $\nu_k \sim \mathcal{N}(0, Q_\nu)$ .

$Q_\omega$ ,  $Q_\nu$  are the **covariance matrices** of the noises: this notation assumes that they are **constant** but in case they are **time-varying**, the following would work in the same way.

# Kalman filter

## The vectorial case

By applying the same reasoning as before, with appropriate algebraic manipulations:

$$\hat{x}_{k+1}^- = A\hat{x}_k + Bu_k, \quad (9a)$$

$$P_{k+1}^- = AP_kA^T + Q_\omega, \quad (9b)$$

$$K_{k+1} = P_{k+1}^- C^T (CP_{k+1}^- C^T + Q_\nu)^{-1}, \quad (9c)$$

$$\hat{x}_{k+1} = \hat{x}_{k+1}^- + K_{k+1}(z_{k+1} - C\hat{x}_{k+1}^-), \quad (9d)$$

$$P_{k+1} = (I - K_{k+1}C)P_{k+1}^-, \quad (9e)$$

where  $I$  is the identity matrix.

The idea behind this is still the **combination of Gaussian random variables**.

# Extended Kalman Filter

## The nonlinear case

Suppose you have a **nonlinear model**:

$$x_{k+1} = f(x_k, u_k, \omega_k), \quad (10a)$$

$$z_k = h(x_k, \nu_k), \quad (10b)$$

with noises as random variables as before.

In this case, we can resort to a **linearization** to **reapply the same reasoning**, but we lose **convergence properties**.

Nonetheless, the **Extended Kalman Filter** (EKF) is a **widely used** heuristic state estimation technique, achieving **good results** in many practical scenarios.

# Extended Kalman Filter

## The linearization

Consider two subsequent time instants  $k$  and  $k + 1$ .

Given the **current estimate**  $\hat{x}_k$  at time  $k$  and the **prediction**  $\hat{x}_{k+1}^-$  at time  $k + 1$ , we can write:

$$f(x_k, u_k, \omega_k) = f(\hat{x}_k, u_k, 0) + F_k(x_k - \hat{x}_k) + W_k\omega_k + o(\|x_k - \hat{x}_k\|), \quad (11a)$$

$$h(x_{k+1}, \nu_{k+1}) = h(\hat{x}_{k+1}^-, 0) + H_{k+1}(x_{k+1} - \hat{x}_{k+1}^-) + L_{k+1}\nu_{k+1} + o(\|x_{k+1} - \hat{x}_{k+1}^-\|), \quad (11b)$$

where  $F = \partial f / \partial x$ ,  $W = \partial f / \partial \omega$ ,  $H = \partial h / \partial x$ , and  $L = \partial h / \partial \nu$  are Jacobian matrices, which we evaluate at appropriate times.

Taking (11a)-(11b) as exact and rearranging known terms, these lead to:

$$x_{k+1} = F_k x_k + f_0(\hat{x}_k, u_k) + W_k \omega_k, \quad (12a)$$

$$z_{k+1} = H_{k+1} x_{k+1} + h_0(\hat{x}_{k+1}^-) + L_{k+1} \nu_{k+1}, \quad (12b)$$

which is easier to manage, both in terms of algebra and **computational complexity**.

# Extended Kalman Filter

## The equations

By applying the same reasoning as before, with appropriate algebraic manipulations:

$$\hat{x}_{k+1}^- = f(\hat{x}_k, u_k, 0), \quad (13a)$$

$$P_{k+1}^- = F_k P_k F_k^T + W_k Q_\omega W_k^T, \quad (13b)$$

$$K_{k+1} = P_{k+1}^- H_{k+1}^T (H_{k+1} P_{k+1}^- H_{k+1}^T + L_{k+1} Q_\nu L_{k+1}^T)^{-1}, \quad (13c)$$

$$\hat{x}_{k+1} = \hat{x}_{k+1}^- + K_{k+1} (z_{k+1} - h(\hat{x}_{k+1}^-, 0)), \quad (13d)$$

$$P_{k+1} = (I - K_{k+1} H_{k+1}) P_{k+1}^-. \quad (13e)$$

It is a **good compromise** between **accuracy** and **computational complexity**.

# Beyond EKF

## The art of filtering

The Kalman Filter and the Extended Kalman Filter gave birth to an entire **family of algorithms** for **state estimation**, **sensor fusion**, and ultimately **localization**.

Some notable examples among the many ones:

- **Unscented Kalman Filter (UKF)**, where the measurement function is not linearized;

# Beyond EKF

## The art of filtering

The Kalman Filter and the Extended Kalman Filter gave birth to an entire **family of algorithms** for **state estimation**, **sensor fusion**, and ultimately **localization**.

Some notable examples among the many ones:

- **Unscented Kalman Filter (UKF)**, where the measurement function is not linearized;
- **Particle Filter (PF)**, where the state is represented by a set of particles each with a likelihood weight, based on Monte Carlo methods;



# Beyond EKF

## The art of filtering

The Kalman Filter and the Extended Kalman Filter gave birth to an entire **family of algorithms** for **state estimation**, **sensor fusion**, and ultimately **localization**.

Some notable examples among the many ones:

- **Unscented Kalman Filter (UKF)**, where the measurement function is not linearized;
- **Particle Filter (PF)**, where the state is represented by a set of particles each with a likelihood weight, based on Monte Carlo methods;
- **Multi-Hypotheses Kalman Filter (MHKF)**, where the number of hypotheses may change over time;

# Beyond EKF

## The art of filtering

The Kalman Filter and the Extended Kalman Filter gave birth to an entire **family of algorithms** for **state estimation**, **sensor fusion**, and ultimately **localization**.

Some notable examples among the many ones:

- **Unscented Kalman Filter (UKF)**, where the measurement function is not linearized;
- **Particle Filter (PF)**, where the state is represented by a set of particles each with a likelihood weight, based on Monte Carlo methods;
- **Multi-Hypotheses Kalman Filter (MHKF)**, where the number of hypotheses may change over time;
- **Switching Kalman Filter (SKF)**, where the system may switch between different models, useful to estimate the state of a hybrid system;

# Beyond EKF

## The art of filtering

The Kalman Filter and the Extended Kalman Filter gave birth to an entire **family of algorithms** for **state estimation**, **sensor fusion**, and ultimately **localization**.

Some notable examples among the many ones:

- **Unscented Kalman Filter (UKF)**, where the measurement function is not linearized;
- **Particle Filter (PF)**, where the state is represented by a set of particles each with a likelihood weight, based on Monte Carlo methods;
- **Multi-Hypotheses Kalman Filter (MHKF)**, where the number of hypotheses may change over time;
- **Switching Kalman Filter (SKF)**, where the system may switch between different models, useful to estimate the state of a hybrid system;
- **Moving Horizon Estimation (MHE)**, where the state is estimated by minimizing a cost function over a finite time horizon.

# robot\_localization

An EKF for ROS 2

One of the very first packages developed for ROS, and one of the first [community projects](#).

It is a flexible and powerful **EKF implementation inside a ROS 2 node**, featuring:

- support for a wide range of **sensor types** using **standard interfaces**;
- **multi-rate** sensor fusion algorithm, with customizable publication rate;
- automatic **sensor data frame conversion** using tf2;
- real-time **tf2 broadcasting**;
- support for both **local** and **global data** with **two chained instances**.

# robot\_localization

## Global localization

When both local (*i.e.*, w.r.t. the starting point) and global (*i.e.*, w.r.t. the world origin) localization data is available, `robot_localization` can be used to **fuse** them, **compensating odometry drift**.

**Two instances** of the node must be active: a **local** one and a **global** one.

**Local:** works in the local **odom frame**, fusing **odometry** and other **inertial data** expressed in body frame; publishes the local pose and tf.

**Global:** works in the global **world frame**, fusing **the same data as the local instance plus global localization data** (e.g., GPS, LiDAR, SLAM); publishes the global pose and tf, and **broadcasts the world → odom tf**, allowing to:

- correct all local globalization data **without resetting the state of the local systems**;
- transform all **data expressed in the local frame** (e.g., point clouds generated by tracking cameras) **compensating odometry drift**.

# robot\_localization

Global localization example: visual odometry and Visual SLAM

# Roadmap

- 1 The perception problem
- 2 Common interfaces
- 3 The tf2 library
- 4 The Extended Kalman Filter
- 5 The mapping problem**
- 6 Simultaneous Localization And Mapping

# The mapping problem

## Definition

Using **exteroceptive sensors**, a robot can gather information about the environment, which can be used to build a **map** of it.

A **map** is a **representation** of the environment, in a format that the robot can **understand**, **parse**, and **store**.

The ultimate goal of mapping is twofold:

- to enable the robot to **localize itself** within the environment;
- to enable **safe navigation** of the robot within the environment.



# The mapping problem

## Challenges

Given the utility requirement of a map, the mapping problem must be **continuously solved in real time**.

Thus, it is challenging because:

- routines must be **efficient**, and run at a sufficiently **high rate**;
- the map must be **accurate**, and **reliable**;
- the map must be in a format that is as much **easy to load and parse** as possible, taking up as little **memory** as possible;
- the map must stay **up-to-date**, and **consistent** with the environment.

# The mapping problem

## Tools for the job

The most important tool for the mapping problem is the **occupancy grid**, a representation of the environment as a **grid** of **cells**, each of which is **occupied** or **free**.

The occupancy grid is a **probabilistic** representation, where each cell is associated with a **probability** of being occupied or free.

The occupancy grid is usually built using **LiDAR** or **camera** depth data, and is updated in real time as the robot moves.

The occupancy grid is the most common representation for **local** and **global** maps.

To efficiently store an occupancy grid, **tree-like data structures** are often employed (e.g., **octrees**).

# The mapping problem

## Tools for the job

The second most important class of tools are **navigation algorithms**, which use the map to plan a **safe** and **efficient** path for the robot to follow.

The definition of such algorithms involves **geometry**, as well as **optimization** and **search** techniques.

They usually rely on two mathematical subjects:

- **topology**, to define the **connectivity** of the map (e.g., Voronoi tessellation);
- **graph theory**, to define the **best way** of moving from one free cell to another (e.g., Dijkstra's, A<sup>\*</sup> algorithms).

# The mapping problem

Tools for the job




Figure 8: Mapping and navigation algorithms execution.

# Roadmap

- 1 The perception problem
- 2 Common interfaces
- 3 The tf2 library
- 4 The Extended Kalman Filter
- 5 The mapping problem
- 6 Simultaneous Localization And Mapping**

# Simultaneous Localization And Mapping

## Definition

The **SLAM** problem is a **chicken-and-egg** problem: a robot must **localize itself** within an environment, while **mapping** it.

The ultimate goal of SLAM is to enable the robot to **navigate** within the environment, while **updating** the map as it moves.

The robot must be able to **efficiently** and **accurately** build a map of the environment, while **localizing itself** within it, with respect to either the origin of the map or the starting point of the robot itself (*i.e.*, with either some or none **prior notion of the environment**).

To solve this problem, **sensor fusion** techniques are often employed, mixing data coming from **heterogeneous sensors** and accounting for **sensor faults**.

Typically, SLAM algorithms rely on recognizable **features** of the environment to build the map and detect motion.

# Simultaneous Localization And Mapping

## Loop closure

While a SLAM system builds a **map** of the environment, it can detect whether it is exploring a zone that it has **already visited**.

When this happens, the system can **close the loop** by **matching** the current "view" with a **previous one**, thus **correcting** the map and the robot's pose.

This process is called **loop closure** and is a key feature of SLAM algorithms.

Loop **detection** and **closing** must also be performed in real time, as well as the subsequent **corrections**; efficient optimization algorithms and data structures are crucial.




Figure 9: Loop closure of the ORB-SLAM2 algorithm.

# Simultaneous Localization And Mapping

Tools for the job

Sensors:

- **LIDARs** for direct environment mapping through depth information;
- **Cameras** to infer the environment structure through image processing;
- **IMUs** to account for the robot's motion and correct sampled data;
- **GNSS** to have a slow, but reliable global position estimate.

Plus all the **algorithms** and the **mathematical tools** we discussed in the context of **mapping**.



# Simultaneous Localization And Mapping

Tools for the job




Figure 10: 3D point cloud generated by a stereoscopic camera.

# 2D SLAM

## Cartographer


**Cartographer** is a **2D** SLAM algorithm developed by Google.

Meant for **offline** floor plan generation, it has been ported in **ROS** for mobile robot navigation.

It used **LiDAR** laser scans, corrected by **IMU** data, to build 2D **submaps** of the **surrounding** environment.

Such maps are then used to **build and update** a **global map**, gluing submaps together and optimizing the result.

Features are particular **environment structures** that are used to detect loop closures and estimate the robot's trajectory.



**Figure 11:** Execution of the Cartographer SLAM algorithm.

# Visual SLAM

## ORB-SLAM

**ORB-SLAM** is a **visual** SLAM algorithm that uses **monocular** or **stereo** cameras to build a map of the environment.

It relies on **binary ORB features** to detect and match points in the environment, building a **covisibility graph** of **keyframes**: relevant views for pose estimation.

The graph, and thus, the map and the camera pose estimate, are optimized with **bundle adjustment**.

Latest versions also include **IMU** samples in the bundle adjustment cost function.

Figure 12: ORB-SLAM2 algorithm execution on a ROS 2 dataset (bag).

Autonomous drone test flight with ORB-SLAM2-based global localization.